

ASSIGNMENT : 1

NAME : NITISH BALORIA

ROLL NO : 2020a1r125

MovieLens 1M Dataset GroupLens Research provides a number of collections of movie ratings data collected from users of MovieLens in the late 1990s and early 2000s. The data provide movie ratings, movie metadata (genres and year), and demographic data about the users (age, zip code, gender identification, and occupation). Such data is often of interest in the development of recommendation systems based on machine learning algorithms. While we do not explore machine learning techniques in detail in this book, I will show you how to slice and dice datasets like these into the exact form you need. The MovieLens 1M dataset contains 1 million ratings collected from 6,000 users on 4,000 movies. It's spread across three tables: ratings, user information, and movie information. After extracting the data from the ZIP file, we can load each table into a pandas DataFrame object using `pandas.read_table` and perform the following task.

1. Perform null values identification in the given dataset.

```
In [16]: import pandas as pd
```

```
In [3]: df = pd.read_csv('C:/Users/user/OneDrive/Desktop/diabetes.csv')
```

```
In [4]: df
```

```
Out[4]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

```
In [5]: null_values = df.isnull().sum()
print(null_values)
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

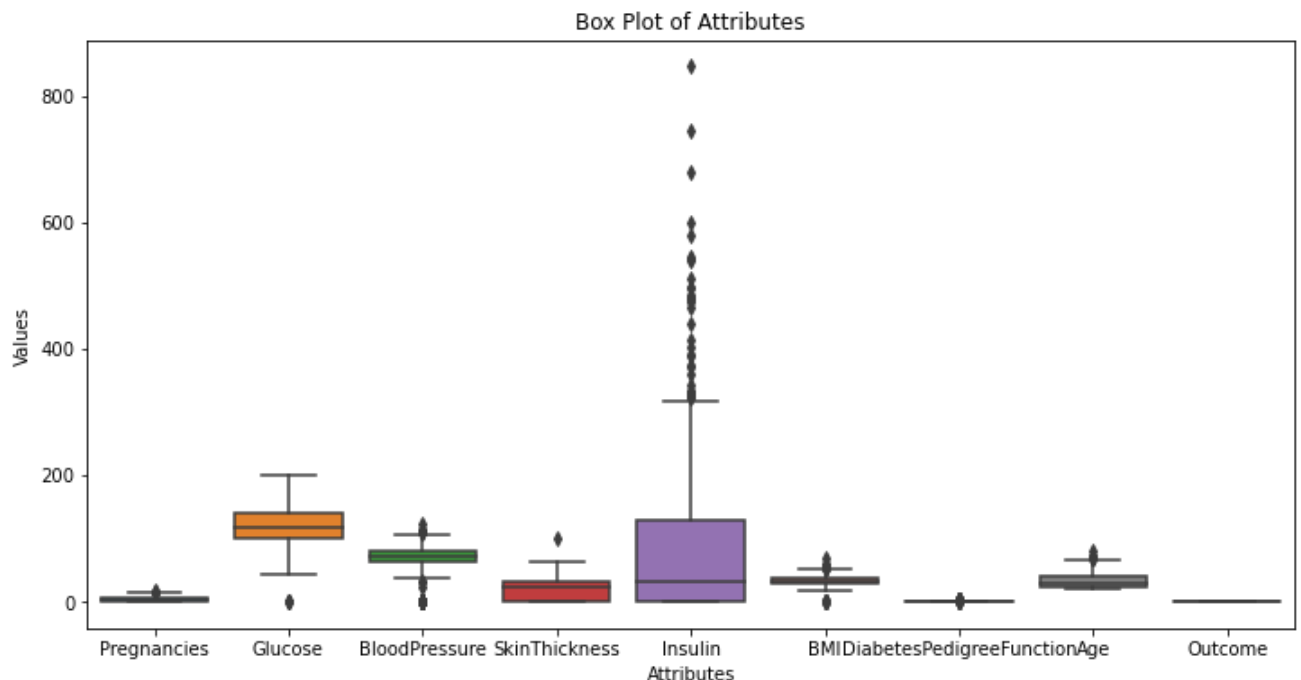
2. Identify types of attributes in the dataset.

```
In [6]: attribute_types = df.dtypes  
print(attribute_types)
```


```
Pregnancies          int64  
Glucose              int64  
BloodPressure        int64  
SkinThickness        int64  
Insulin              int64  
BMI                  float64  
DiabetesPedigreeFunction float64  
Age                  int64  
Outcome              int64  
dtype: object
```

3. Plot Box plot and violin plot. (also state the inference of each attribute and also find the outlier in the attribute)

```
In [8]: import matplotlib.pyplot as plt  
import seaborn as sns  
  
# CREATE BOX PLOTS  
plt.figure(figsize=(12, 6))  
sns.boxplot(data=df)  
plt.title('Box Plot of Attributes')  
plt.xlabel('Attributes')  
plt.ylabel('Values')  
plt.show()
```



```
In [9]: # CREATE VIOLIN PLOTS
plt.figure(figsize=(12, 6))
sns.violinplot(data=df)
plt.title('Violin Plot of Attributes')
plt.xlabel('Attributes')
plt.ylabel('Values')
plt.show()
```



```
In [10]: # IDENTIFY OUTLIER USING BOX PLOTS
for column in df.columns:
    if df[column].dtype != 'object': # Exclude non-numeric attributes
        q1 = df[column].quantile(0.25)
        q3 = df[column].quantile(0.75)
        iqr = q3 - q1
        lower_bound = q1 - 1.5 * iqr
        upper_bound = q3 + 1.5 * iqr
        outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
        if outliers.shape[0] > 0:
            print(f"Outliers in attribute '{column}':")
            print(outliers)
```

```
502      DiabetesPedigreeFunction  Age  Outcome
75                0.140    22         0
182                0.299    21         0
342                0.389    22         0
349                0.346    37         1
502                0.727    41         1
Outliers in attribute 'BloodPressure':
   Pregnancies  Glucose  BloodPressure  SkinThickness  Insulin  BMI  \
7             10     115             0              0        0  35.3
15             7      100             0              0        0  30.0
18             1     103             30             38       83  43.3
43             9     171            110             24     240  45.4
49             7     105              0              0        0   0.0
60             2       84              0              0        0   0.0
78             0     131              0              0        0  43.2
81             2       74              0              0        0   0.0
84             5     137            108              0        0  48.8
106            1       96            122              0        0  22.4
```

```
In [11]: # IDENTIFY INFERENCE FOR EACH ATTRIBUTE
for column in df.columns:
    if df[column].dtype == 'object': # CATEGORICAL ATTRIBUTES
        print(f"Attribute '{column}' is a categorical attribute.")
        print(f"Unique values: {df[column].unique()}")
    else: # NUMERIC ATTRIBUTES
        print(f"Attribute '{column}' is a numeric attribute.")
        print(f"Summary statistics: {df[column].describe()}")
```

Attribute 'Pregnancies' is a numeric attribute.

Summary statistics: count 768.000000

mean 3.845052

std 3.369578

min 0.000000

25% 1.000000

50% 3.000000

75% 6.000000

max 17.000000

Name: Pregnancies, dtype: float64

Attribute 'Glucose' is a numeric attribute.

Summary statistics: count 768.000000

mean 120.894531

std 31.972618

min 0.000000

25% 99.000000

50% 117.000000

75% 140.250000

max 199.000000

Name: Glucose, dtype: float64

Attribute 'BloodPressure' is a numeric attribute.

Summary statistics: count 768.000000

mean 69.105469

std 19.355807

min 0.000000

25% 62.000000

50% 72.000000

75% 80.000000

max 122.000000

Name: BloodPressure, dtype: float64

Attribute 'SkinThickness' is a numeric attribute.

Summary statistics: count 768.000000

mean 20.536458

std 15.952218

min 0.000000

25% 0.000000

50% 23.000000

75% 32.000000

max 99.000000

Name: SkinThickness, dtype: float64

Attribute 'Insulin' is a numeric attribute.

Summary statistics: count 768.000000

mean 79.799479

std 115.244002

min 0.000000

25% 0.000000

50% 30.500000

75% 127.250000

max 846.000000

Name: Insulin, dtype: float64

Attribute 'BMI' is a numeric attribute.

Summary statistics: count 768.000000

mean 31.992578

std 7.884160

min 0.000000

25% 27.300000

50% 32.000000

75% 36.600000

max 67.100000

Name: BMI, dtype: float64

Attribute 'DiabetesPedigreeFunction' is a numeric attribute.

Summary statistics: count 768.000000

mean 0.471876

std 0.331329

min 0.078000

25% 0.243750

50% 0.372500

75% 0.626250

max 2.420000

Name: DiabetesPedigreeFunction, dtype: float64

Attribute 'Age' is a numeric attribute.

Summary statistics: count 768.000000

```

mean      33.240885
std       11.760232
min       21.000000
25%       24.000000
50%       29.000000
75%       41.000000
max       81.000000
Name: Age, dtype: float64
Attribute 'Outcome' is a numeric attribute.
Summary statistics: count    768.000000
mean      0.348958
std       0.476951
min       0.000000
25%       0.000000
50%       0.000000
75%       1.000000
max       1.000000
Name: Outcome, dtype: float64

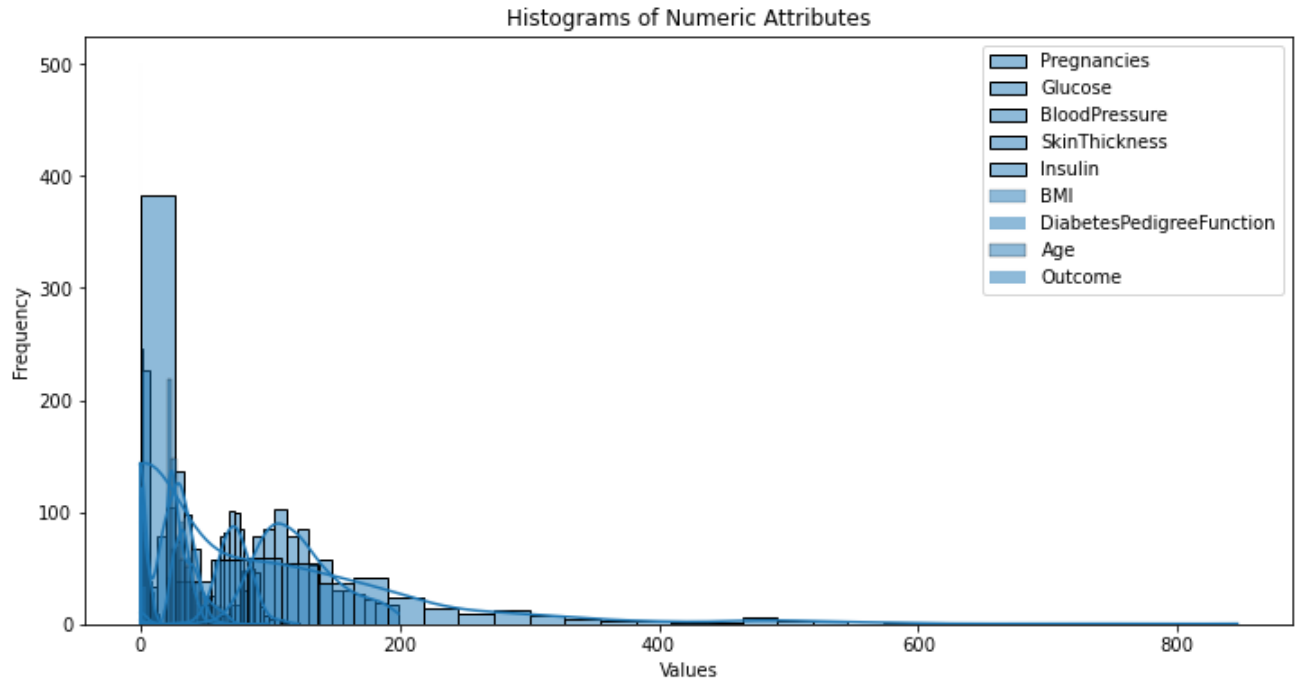
```

4. Histogram and identification of overlapping.(also state the inference for each attribute.)

```

In [14]: # CREATE HISTOGRAM
plt.figure(figsize=(12, 6))
for column in df.columns:
    if df[column].dtype != 'object': # Exclude non-numeric attributes
        sns.histplot(df[column], kde=True, label=column)
plt.title('Histograms of Numeric Attributes')
plt.xlabel('Values')
plt.ylabel('Frequency')
plt.legend()
plt.show()

```



```
In [13]: # IDENTIFY OVERLAPPING
for column1 in df.columns:
    for column2 in df.columns:
        if column1 != column2 and df[column1].dtype != 'object' and df[column2].dtype != 'object':
            overlap = df[(df[column1] > df[column2].min()) & (df[column1] < df[column2].max())]
            if overlap.shape[0] > 0:
                print(f"Overlapping between '{column1}' and '{column2}':")
                print(overlap)
```

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
5	0.201	30	0
..
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[657 rows x 9 columns]

Overlapping between 'Pregnancies' and 'DiabetesPedigreeFunction':

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
1	1	85	66	29	0	26.6	
3	1	89	66	23	94	28.1	
8	2	197	70	45	543	30.5	

```
In [15]: # IDENTIFY INFERENCE FOR EACH ATTRIBUTE
for column in df.columns:
    if df[column].dtype == 'object': # CATEGORICAL ATTRIBUTES
        print(f"Attribute '{column}' is a categorical attribute.")
        print(f"Unique values: {df[column].unique()}")
    else: # NUMERIC ATTRIBUTES
        print(f"Attribute '{column}' is a numeric attribute.")
        print(f"Summary statistics: {df[column].describe()}")
```


Attribute 'Pregnancies' is a numeric attribute.

Summary statistics: count 768.000000

mean 3.845052

std 3.369578

min 0.000000

25% 1.000000

50% 3.000000

75% 6.000000

max 17.000000

Name: Pregnancies, dtype: float64

Attribute 'Glucose' is a numeric attribute.

Summary statistics: count 768.000000

mean 120.894531

std 31.972618

min 0.000000

25% 99.000000

50% 117.000000

75% 140.250000

max 199.000000

Name: Glucose, dtype: float64

Attribute 'BloodPressure' is a numeric attribute.

Summary statistics: count 768.000000

mean 69.105469

std 19.355807

min 0.000000

25% 62.000000

50% 72.000000

75% 80.000000

max 122.000000

Name: BloodPressure, dtype: float64

Attribute 'SkinThickness' is a numeric attribute.

Summary statistics: count 768.000000

mean 20.536458

std 15.952218

min 0.000000

25% 0.000000

50% 23.000000

75% 32.000000

max 99.000000

Name: SkinThickness, dtype: float64

Attribute 'Insulin' is a numeric attribute.

Summary statistics: count 768.000000

mean 79.799479

std 115.244002

min 0.000000

25% 0.000000

50% 30.500000

75% 127.250000

max 846.000000

Name: Insulin, dtype: float64

Attribute 'BMI' is a numeric attribute.

Summary statistics: count 768.000000

mean 31.992578

std 7.884160

min 0.000000

25% 27.300000

50% 32.000000

75% 36.600000

max 67.100000

Name: BMI, dtype: float64

Attribute 'DiabetesPedigreeFunction' is a numeric attribute.

Summary statistics: count 768.000000

mean 0.471876

std 0.331329

min 0.078000

25% 0.243750

50% 0.372500

75% 0.626250

max 2.420000

Name: DiabetesPedigreeFunction, dtype: float64

Attribute 'Age' is a numeric attribute.

Summary statistics: count 768.000000

```

mean      33.240885
std       11.760232
min       21.000000
25%      24.000000
50%      29.000000
75%      41.000000
max       81.000000
Name: Age, dtype: float64
Attribute 'Outcome' is a numeric attribute.
Summary statistics: count      768.000000
mean      0.348958
std       0.476951
min       0.000000
25%      0.000000
50%      0.000000
75%      1.000000
max       1.000000
Name: Outcome, dtype: float64

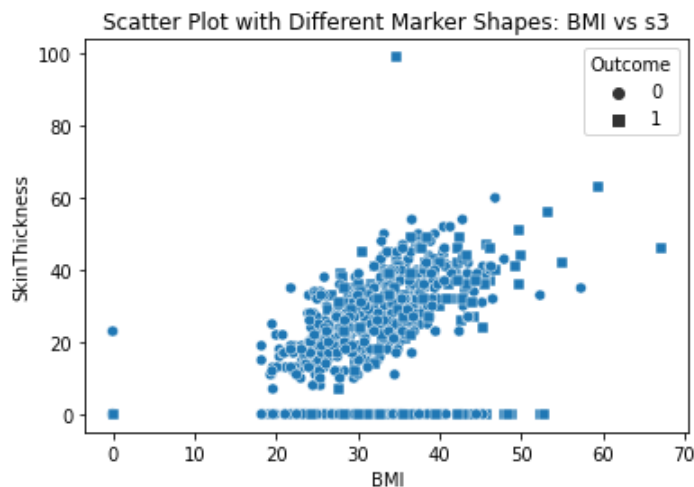
```

5. Draw different types of scatter plot.(using seaborn library)

```

In [29]: # SCATTER PLOT WITH SINGLE VARIABLES(UNIVARIATE SCATTER PLOT)
sns.scatterplot(x='Age', y='BMI', data=df)
plt.title('Scatter Plot: Age vs BMI')
plt.xlabel('Age')
plt.ylabel('BMI')
plt.show()

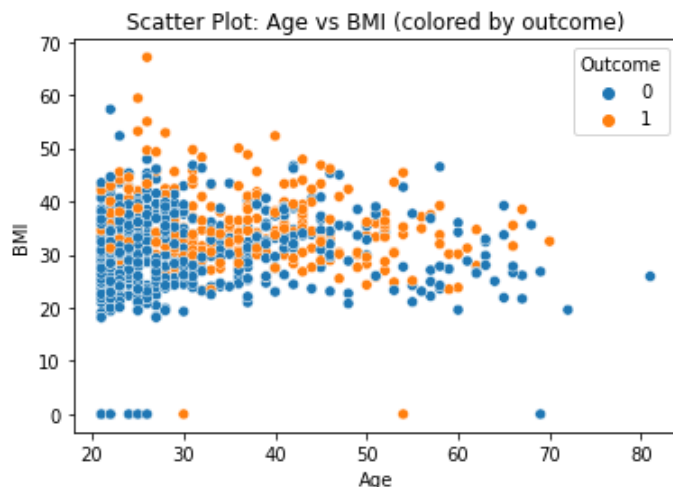
```



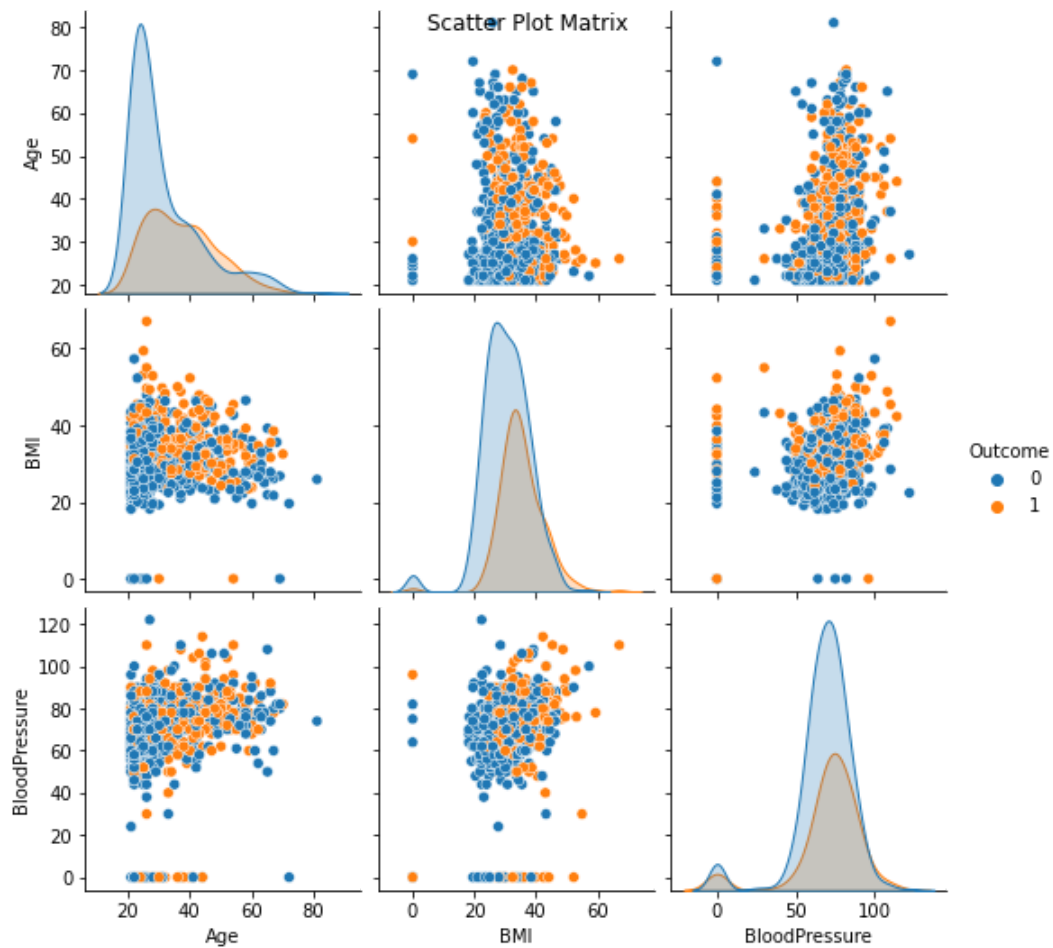
```

In [30]: # SCATTER PLOT WITH TWO VARIABLES(BIVARIATE SCATTER PLOT)
sns.scatterplot(x='Age', y='BMI', hue='Outcome', data=df)
plt.title('Scatter Plot: Age vs BMI (colored by outcome)')
plt.xlabel('Age')
plt.ylabel('BMI')
plt.show()


```




```
In [31]: # SCATTER PLOT MATRIX(PAIRPLOT)
sns.pairplot(data=df, vars=['Age', 'BMI', 'BloodPressure'], hue='Outcome')
plt.suptitle('Scatter Plot Matrix')
plt.show()
```



```
In [32]: # SCATTER PLOT WITH REGRESSION LINE(REGRESSION PLOT)
sns.regplot(x='BMI', y='BloodPressure', data=df)
plt.title('Scatter Plot with Regression Line: BMI vs bp')
plt.xlabel('BMI')
plt.ylabel('BloodPressure')
plt.show()
```




```
In [33]: # SCATTER PLOT WITH DIFFERENT MARKER SHAPES(SCATTER PLOT WITH MARKER PARAMETER)
sns.scatterplot(x='BMI', y='SkinThickness', data=df, style='Outcome', markers=['o', 's'])
plt.title('Scatter Plot with Different Marker Shapes: BMI vs SkinThickness')
plt.xlabel('BMI')
plt.ylabel('SkinThickness')
plt.show()
```




6. Univariate and multivariate analysis.

```
In [41]: # UNIVARIATE ANALYSIS: HISTOGRAM
sns.histplot(df['Age'], bins=10, kde=True)
plt.title('Histogram: Age')
plt.xlabel('Age')
plt.ylabel('Frequency')
plt.show()
```



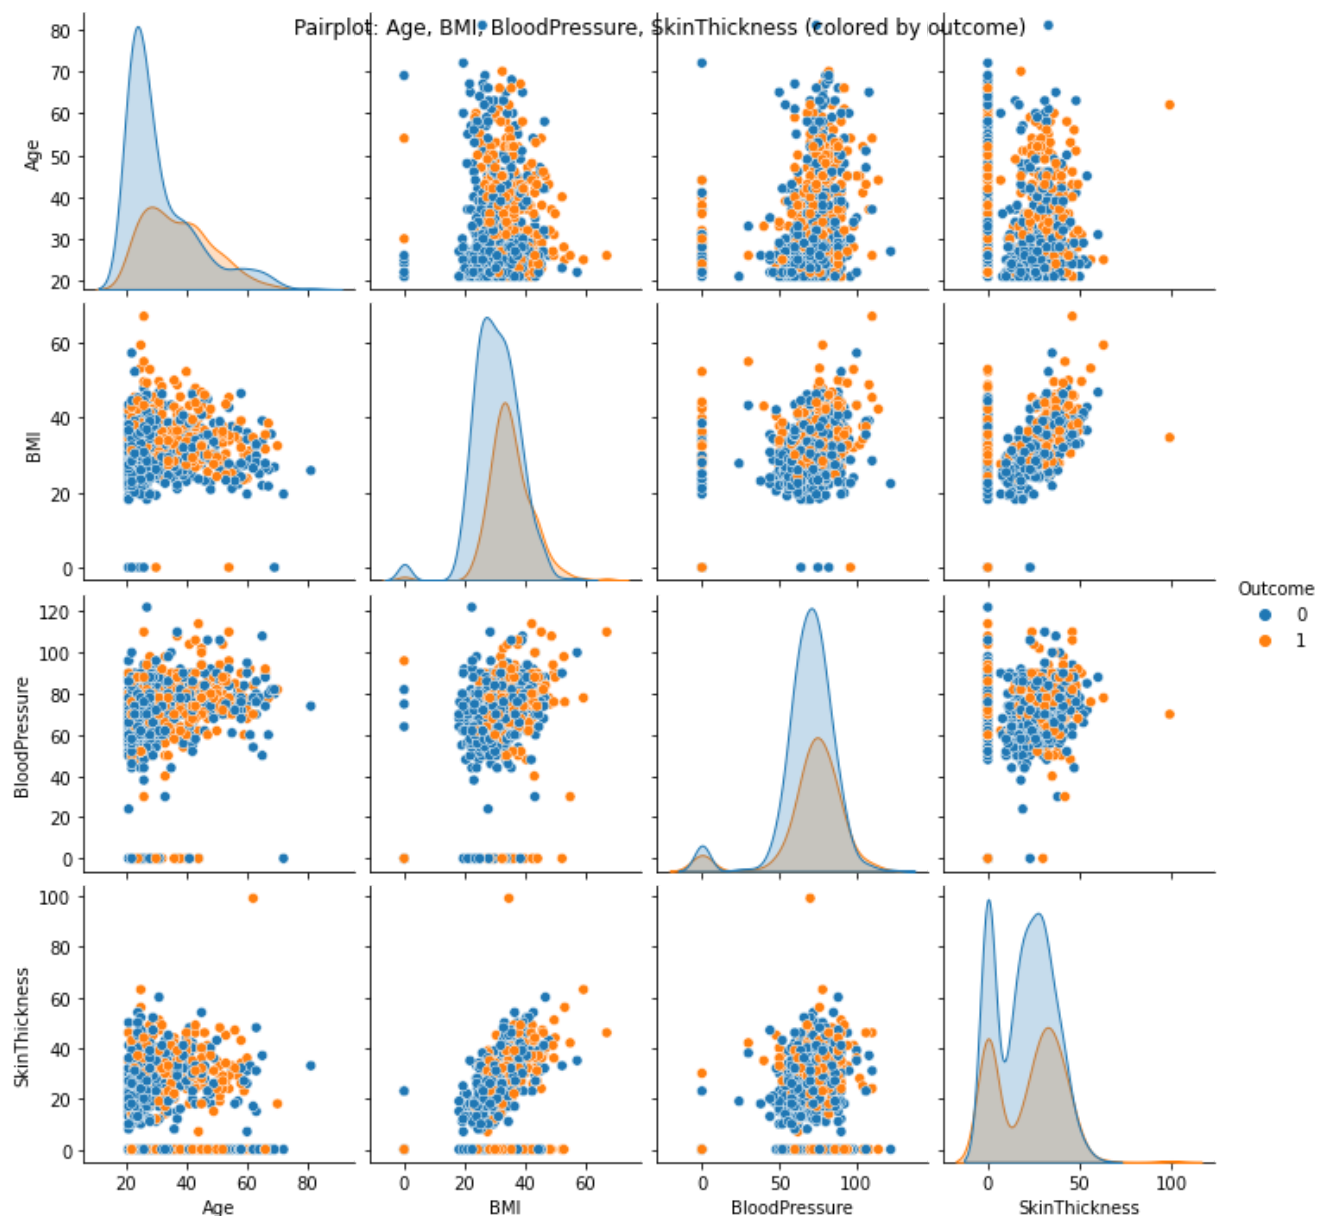
```
In [40]: # UNIVARIATE ANALYSIS: BOX PLOT
```

```
sns.boxplot(x='Age', data=df)
plt.title('Box Plot: Age')
plt.xlabel('Age')
plt.show()
```




```
In [42]: # MULTIVARIATE ANALYSIS: PAIRPLOT
```

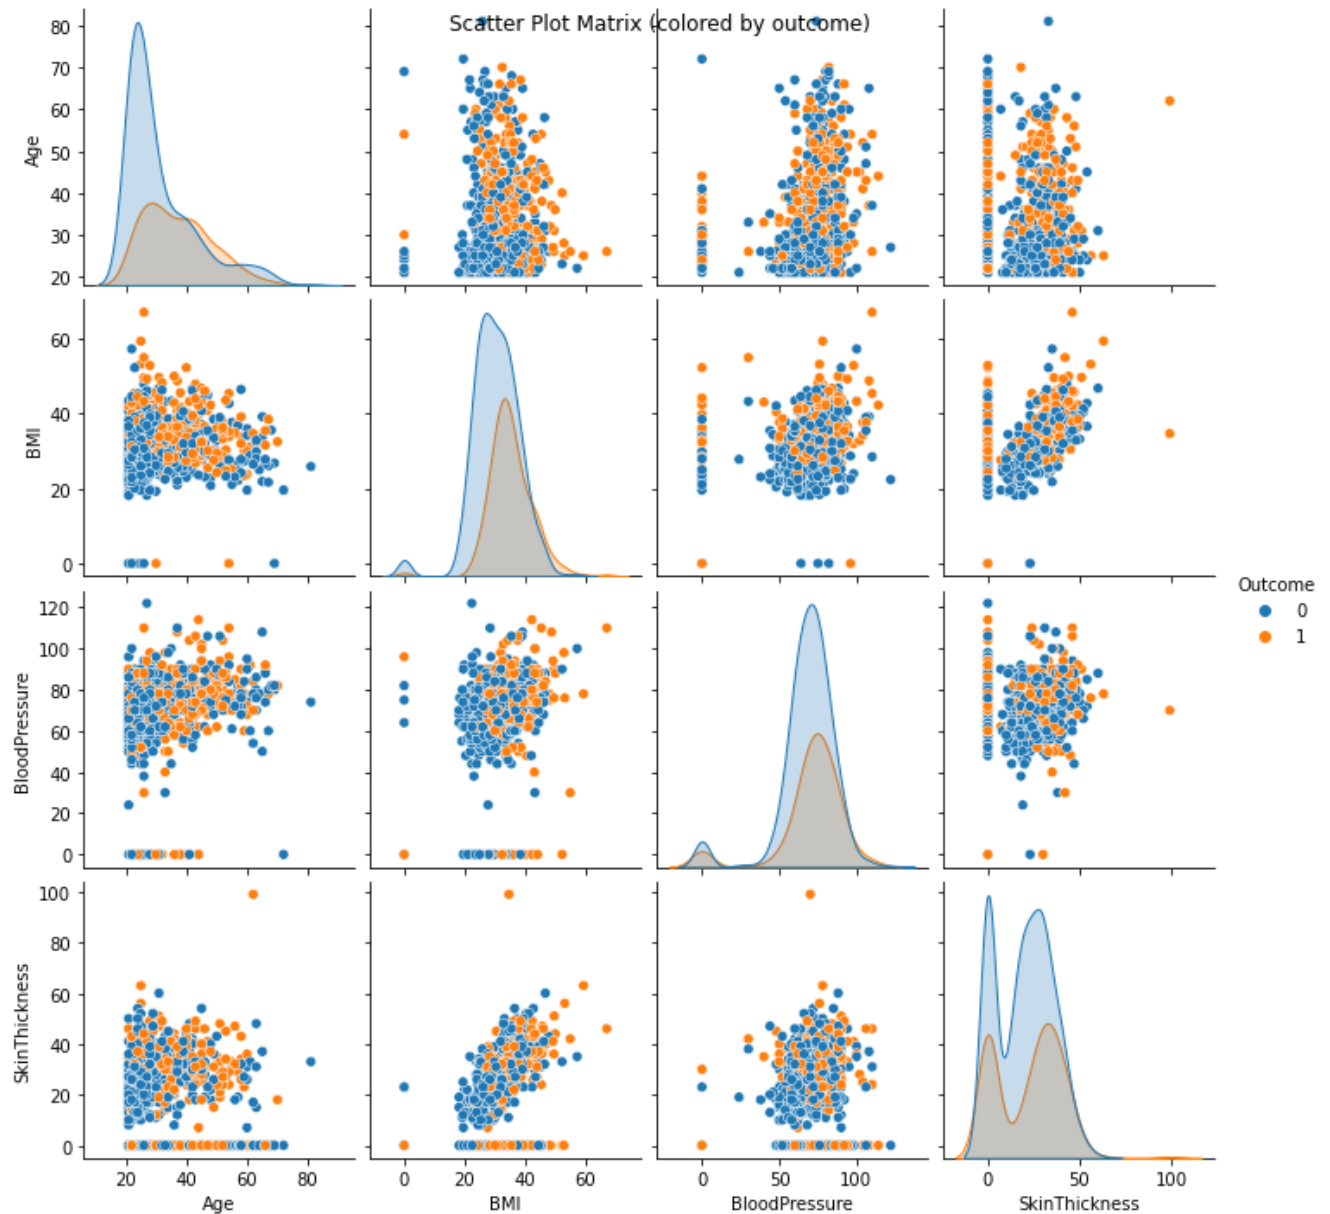
```
sns.pairplot(data=df, hue='Outcome', vars=['Age', 'BMI', 'BloodPressure', 'SkinThickness'])
plt.suptitle('Pairplot: Age, BMI, BloodPressure, SkinThickness (colored by outcome)')
plt.show()
```



```
In [43]: # MULTIVARIATE ANALYSIS: HEATMAP
corr_matrix = df.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Heatmap: Correlation between variables')
plt.show()
```



```
In [44]: # MULTIVARIATE ANALYSIS: SCATTER PLOT MATRIX
sns.pairplot(data=df, vars=['Age', 'BMI', 'BloodPressure', 'SkinThickness'], hue='Outcome')
plt.suptitle('Scatter Plot Matrix (colored by outcome)')
plt.show()
```



2. Diabetics datasets :

Data Exploration: This includes inspecting the data, visualizing the data, and cleaning the data. Some of the steps used are as follows:

1. Viewing the data statistics.

```
In [52]: # VIEW THE FIRST FEW ROWS OF THE DATA
print("Head of the data:")
print(df.head())
```

Head of the data:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

	DiabetesPedigreeFunction	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1

```
In [47]: # VIEW THE DATA STATISTICS
print("Data statistics:")
print(df.describe())
```

Data statistics:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	\
count	768.000000	768.000000	768.000000	768.000000	768.000000	
mean	3.845052	120.894531	69.105469	20.536458	79.799479	
std	3.369578	31.972618	19.355807	15.952218	115.244002	
min	0.000000	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	99.000000	62.000000	0.000000	0.000000	
50%	3.000000	117.000000	72.000000	23.000000	30.500000	
75%	6.000000	140.250000	80.000000	32.000000	127.250000	
max	17.000000	199.000000	122.000000	99.000000	846.000000	

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000

```
In [48]: # View the data types of each column
print("Data types:")
print(df.dtypes)
```

Data types:

Pregnancies	int64
Glucose	int64
BloodPressure	int64
SkinThickness	int64
Insulin	int64
BMI	float64
DiabetesPedigreeFunction	float64
Age	int64
Outcome	int64
dtype:	object

```
In [49]: # VIEW THE SHAPE OF THE DATA (NUMBER OF ROWS AND COLUMNS)
print("Data shape:")
print(df.shape)
```

Data shape:
(768, 9)


```
In [50]: # VIEW THE COLUMN NAMES
print("Column names:")
print(df.columns)
```

```
Column names:
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
In [51]: # VIEW THE NUMBER OF MISSING VALUES IN EACH COLUMN
print("Missing values:")
print(df.isnull().sum())
```

```
Missing values:
Pregnancies           0
Glucose               0
BloodPressure         0
SkinThickness         0
Insulin              0
BMI                  0
DiabetesPedigreeFunction  0
Age                  0
Outcome              0
dtype: int64
```

2. Finding out the dimensions of the dataset, the variable names, the data types, etc.

```
In [59]: # VIEW THE SHAPE OF THE DATA (NUMBER OF ROWS AND COLUMNS)
print("Data shape:")
print(df.shape)
```

```
Data shape:
(768, 9)
```

```
In [57]: # VIEW THE COLUMN NAMES
print("Column names:")
print(df.columns)
```

```
Column names:
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')
```

```
In [58]: # VIEW THE DATA TYPES OF EACH COLUMN
print("Data types:")
print(df.dtypes)
```

```
Data types:
Pregnancies           int64
Glucose              int64
BloodPressure         int64
SkinThickness         int64
Insulin              int64
BMI                  float64
DiabetesPedigreeFunction float64
Age                  int64
Outcome              int64
dtype: object
```

3. Checking for null values.


```
In [62]: # CHECK FOR MISSING VALUES IN EACH COLUMN
print("Missing values:")
print(df.isnull().sum())
```

```
Missing values:
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
dtype: int64
```


4. Inspecting the target variable using pie plot and count plot.

```
In [65]: # INSPECT THE TARGET VARIABLE (e.g., 'Outcome') USING A PIE PLOT
plt.figure(figsize=(6, 6))
df['Outcome'].value_counts().plot.pie(labels=['Non-diabetic', 'Diabetic'], autopct='%1.1f%%', sta
plt.title('Distribution of Outcome (Diabetic vs Non-diabetic)')
plt.show()
```

Distribution of Outcome (Diabetic vs Non-diabetic)




```
In [66]: # INSPECT THE TARGET VARIABLE USING A COUNT PLOT
plt.figure(figsize=(6, 6))
sns.countplot(x='Outcome', data=df)
plt.xlabel('Outcome')
plt.ylabel('Count')
plt.title('Distribution of Outcome (Diabetic vs Non-diabetic)')
plt.show()
```



5. Finding out the correlation among different features using heatmap and the bivariate relation between each pair of features using pair plot.

```
In [70]: # COMPUTE THE CORRELATION MATRIX AMONG FEATURES
corr_matrix = df.corr()

# CREATE A HEATMAP TO VISUALIZE THE CORRELATION MATRIX
plt.figure(figsize=(10, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Heatmap of Features')
plt.show()
```



```
In [71]: # CREATE A PAIR PLOT TO VISUALIZE THE BIVARIATE RELATION BETWEEN FEATURES
sns.pairplot(df, diag_kind='kde', hue='Outcome', markers='o', plot_kws={'alpha':0.5})
plt.suptitle('Pair Plot of Features by Outcome', y=1.02)
plt.show()
```



Model Training: 5 Classification Algorithms have been used to find out the best one. These are Logistic Regression, Support Vector Machine, Random Forest, K-Nearest Neighbours, and Naive Bayes.

In each of the algorithms, the steps followed are as follows:

1. Importing the library for the algorithm.

```
In [77]: # Import Libraries for Logistic Regression
from sklearn.linear_model import LogisticRegression
```

```
In [73]: # Import Libraries for Support Vector Machine
from sklearn.svm import SVC
```

```
In [74]: # Import Libraries for Random Forest
from sklearn.ensemble import RandomForestClassifier
```

```
In [75]: # Import Libraries for K-Nearest Neighbors
from sklearn.neighbors import KNeighborsClassifier
```

```
In [76]: # Import Libraries for Naive Bayes
from sklearn.naive_bayes import GaussianNB
```

2. Creating an instance of the Classifier (with default values of parameters or by specifying certain values in certain cases).

```
In [83]: # Create an instance of Logistic Regression with default parameters
logistic_regression = LogisticRegression()
```

```
In [79]: # Create an instance of Support Vector Machine with specified parameters
svm = SVC(kernel='linear', C=1.0)
```

```
In [80]: # Create an instance of Random Forest with default parameters
random_forest = RandomForestClassifier()
```

```
In [81]: # Create an instance of K-Nearest Neighbors with specified parameters
knn = KNeighborsClassifier(n_neighbors=5)
```

```
In [82]: # Create an instance of Naive Bayes with default parameters
naive_bayes = GaussianNB()
```

3. Training the model on the train set.

```
In [96]: X = df.drop('Outcome', axis=1) # Drop the 'Outcome' column as features (X)
y = df['Outcome']

# Split the dataset into training and testing sets
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Create instances of classifiers with default parameter values
logreg = LogisticRegression()
svm = SVC()
rf = RandomForestClassifier()
knn = KNeighborsClassifier()
nb = GaussianNB()

# Train the classifiers on the training set
logreg.fit(X_train, y_train)
svm.fit(X_train, y_train)
rf.fit(X_train, y_train)
knn.fit(X_train, y_train)
nb.fit(X_train, y_train)
```

C:\Users\user\anaconda\lib\site-packages\sklearn\linear_model_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
n_iter_i = _check_optimize_result(

```
Out[96]: GaussianNB()
```

4. Prediction on the test set using the trained model.

```
In [95]: # Import necessary libraries
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score
# Assuming the classifiers (logreg, svm, rf, knn, nb) are already trained on the training set
# Make predictions on the test set
y_pred_logreg = logreg.predict(X_test)
y_pred_svm = svm.predict(X_test)
y_pred_rf = rf.predict(X_test)
y_pred_knn = knn.predict(X_test)
y_pred_nb = nb.predict(X_test)

# Evaluate the performance of the classifiers
print('Logistic Regression Accuracy:', accuracy_score(y_test, y_pred_logreg))
print('Support Vector Machine Accuracy:', accuracy_score(y_test, y_pred_svm))
print('Random Forest Accuracy:', accuracy_score(y_test, y_pred_rf))
print('K-Nearest Neighbors Accuracy:', accuracy_score(y_test, y_pred_knn))
print('Naive Bayes Accuracy:', accuracy_score(y_test, y_pred_nb))
print('---')
print('Logistic Regression Precision:', precision_score(y_test, y_pred_logreg))
print('Support Vector Machine Precision:', precision_score(y_test, y_pred_svm))
print('Random Forest Precision:', precision_score(y_test, y_pred_rf))
print('K-Nearest Neighbors Precision:', precision_score(y_test, y_pred_knn))
print('Naive Bayes Precision:', precision_score(y_test, y_pred_nb))
print('---')
print('Logistic Regression Recall:', recall_score(y_test, y_pred_logreg))
print('Support Vector Machine Recall:', recall_score(y_test, y_pred_svm))
print('Random Forest Recall:', recall_score(y_test, y_pred_rf))
print('K-Nearest Neighbors Recall:', recall_score(y_test, y_pred_knn))
print('Naive Bayes Recall:', recall_score(y_test, y_pred_nb))
print('---')
print('Logistic Regression F1 Score:', f1_score(y_test, y_pred_logreg))
print('Support Vector Machine F1 Score:', f1_score(y_test, y_pred_svm))
print('Random Forest F1 Score:', f1_score(y_test, y_pred_rf))
print('K-Nearest Neighbors F1 Score:', f1_score(y_test, y_pred_knn))
print('Naive Bayes F1 Score:', f1_score(y_test, y_pred_nb))

Logistic Regression Accuracy: 0.7532467532467533
Support Vector Machine Accuracy: 0.7662337662337663
Random Forest Accuracy: 0.7337662337662337
K-Nearest Neighbors Accuracy: 0.6623376623376623
Naive Bayes Accuracy: 0.7662337662337663
---
Logistic Regression Precision: 0.6491228070175439
Support Vector Machine Precision: 0.7209302325581395
Random Forest Precision: 0.625
K-Nearest Neighbors Precision: 0.5245901639344263
Naive Bayes Precision: 0.6610169491525424
---
Logistic Regression Recall: 0.6727272727272727
Support Vector Machine Recall: 0.5636363636363636
Random Forest Recall: 0.6363636363636364
K-Nearest Neighbors Recall: 0.5818181818181818
Naive Bayes Recall: 0.7090909090909091
---
Logistic Regression F1 Score: 0.6607142857142858
Support Vector Machine F1 Score: 0.6326530612244898
Random Forest F1 Score: 0.6306306306306306
K-Nearest Neighbors F1 Score: 0.5517241379310346
Naive Bayes F1 Score: 0.6842105263157895
```

5. Calculating the accuracy of the prediction.

```
In [97]: from sklearn.metrics import accuracy_score

# Assuming the classifiers (logreg, svm, rf, knn, nb) are already trained on the training set
# Assuming the predictions (y_pred_logreg, y_pred_svm, y_pred_rf, y_pred_knn, y_pred_nb) are already calculated

# Calculate accuracy for each classifier
accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
accuracy_svm = accuracy_score(y_test, y_pred_svm)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
accuracy_knn = accuracy_score(y_test, y_pred_knn)
accuracy_nb = accuracy_score(y_test, y_pred_nb)

# Print accuracy for each classifier
print('Logistic Regression Accuracy:', accuracy_logreg)
print('Support Vector Machine Accuracy:', accuracy_svm)
print('Random Forest Accuracy:', accuracy_rf)
print('K-Nearest Neighbors Accuracy:', accuracy_knn)
print('Naive Bayes Accuracy:', accuracy_nb)
```

```
Logistic Regression Accuracy: 0.7532467532467533
Support Vector Machine Accuracy: 0.7662337662337663
Random Forest Accuracy: 0.7337662337662337
K-Nearest Neighbors Accuracy: 0.6623376623376623
Naive Bayes Accuracy: 0.7662337662337663
```